

# Tutorial 8 - Belief Propagation and Sampling

## COMP9418 – Advanced Topics in Statistical Machine Learning

Lecturer: Gustavo Batista

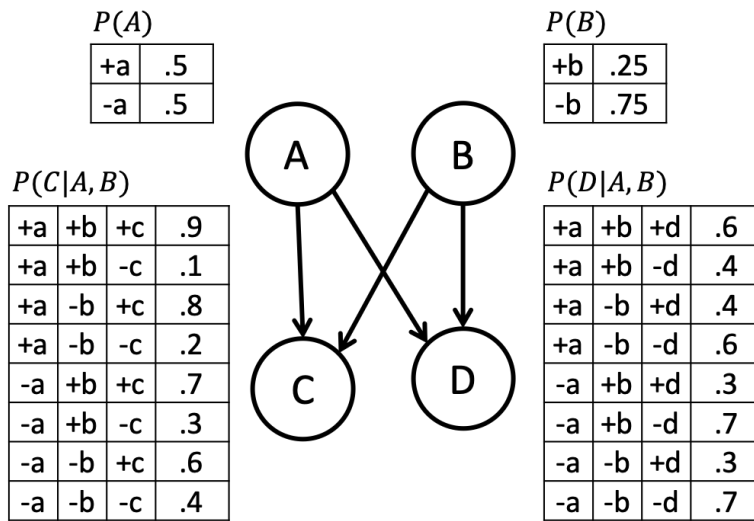
**Lecture:** Belief Propagation and Sampling

**Topic:** Questions from lecture topics

**Last revision:** Thursday 19<sup>th</sup> November, 2020 at 15:21

### Question 1

Consider the following Bayesian network:



Suppose we condition on evidence  $e : D = \text{true}$ . Suppose that we have run the Parallel Iterative Belief Propagation (IBP) algorithm on the network, where it converges and yields the following set of messages and family marginals:

$A$	$\pi_C(A)$	$\pi_D(A)$	$\lambda_C(A)$	$\lambda_D(A)$
+a		.5	.5	.6
-a		.5	.5	.4

$B$	$\pi_C(B)$	$\pi_D(B)$	$\lambda_C(B)$	$\lambda_D(B)$
+b	.3	.25	.5	
-b	.7	.75	.5	

$A$	$B$	$C$	$\beta(A, B, C)$
+a	+b	+c	.162
+a	+b	-c	.018
+a	-b	+c	
+a	-b	-c	.084
-a	+b	+c	.084
-a	+b	-c	.036
-a	-b	+c	.168
-a	-b	-c	.112

$A$	$B$	$D$	$\beta(A, B, D)$
+a	+b	+d	.2
+a	+b	-d	0
+a	-b	+d	
+a	-b	-d	0
-a	+b	+d	.1
-a	+b	-d	0
-a	-b	+d	.3
-a	-b	-d	0

- Fill in the missing values for IBP messages.
- Fill in the missing values for family marginals.
- Compute marginals  $\beta(A)$  and  $\beta(B)$  using the IBP messages and those computed in (a).
- Compute marginals  $\beta(A)$  and  $\beta(B)$  by summing out the appropriate variables from the family marginals  $\beta(ABC)$  as well as  $\beta(ABD)$ .
- Compute joint marginal  $\beta(AB)$  by summing out the appropriate variables from family marginals  $\beta(ABC)$  as well as  $\beta(ABD)$ .
- Are the marginals computed in (d) consistent? What about those computed in (e)?

Consider the following IBP algorithm:


```

Data:  $N$ : Bayesian network
Data:  $e$ : evidence
Result: Approximate marginals,  $\beta(XU)$  of  $P(XU|e)$  for each family  $XU$  in  $N$ 
1 begin
2    $t \leftarrow 0$ ;
3   initialize all messages;
4   while messages have not converged do
5      $t \leftarrow t + 1$ ;
6     for each node  $X$  with parents  $U$  do
7       for each parent  $U_i$  do
8          $\lambda_X^t(U_i) = \eta \sum_{X \setminus \{U_i\}} \lambda_e(X) \phi_X(X, U) \prod_{k \neq i} \pi_X^{t-1}(U_k) \prod_j \lambda_{Y_j}^{t-1}(X)$ ;
9       end
10      for each child  $Y_j$  do
11         $\pi_{Y_j}^t(X) = \eta \sum_U \lambda_e(X) \phi_X(X, U) \prod_i \pi_X^{t-1}(U_i) \prod_{k \neq j} \lambda_{Y_k}^{t-1}(X)$ ;
12      end
13    end
14  end
15 end
16 return  $\beta(XU) = \eta \lambda_e(X) \phi_X(X, U) \prod_i \pi_X^t(U_i) \prod_j \lambda_{Y_j}^t(X)$ 

```

**Answer**

a. The following figure shows all messages:



Let's start computing  $\pi_C(A)$ . According to the IBP algorithm:

$$\pi_C(A) = \eta \phi_A(A) \lambda_D(A)$$

This simplified equation comes from the fact that node  $A$  has no parents and no evidence. Also, this message does not consider the message  $\lambda_C(A)$  coming from node  $C$  because we are computing a message directed to the same node  $C$ .

Therefore,

$A$	$\pi_C(A)$
+a	.6
-a	.4

Remember that the multiplication by the constant  $\eta$  renormalizes the message  $\pi_C(A)$  to sum to one.

Now, we compute  $\lambda_D(B)$  using the equation from the IBP algorithm:

$$\lambda_D(B) = \eta \sum_{A,B} \lambda_e(D) \phi_D(D, A, B) \pi_D(A)$$

Also, notice that node  $B$  has no children, and we should not consider the message  $\pi_D(B)$  when computing a message to  $B$ .

We first compute the multiplication  $\lambda_e(D) \phi_D(D, A, B) \pi_D(A)$ .

$A$	$B$	$D$	$\lambda_e(D) \phi_D(D, A, B) \pi_D(A)$
+a	+b	+d	.3
+a	+b	-d	0
+a	-b	+d	.2
+a	-b	-d	0
-a	+b	+d	.15
-a	+b	-d	0
-a	-b	+d	.15
-a	-b	-d	0

After the elimination of variables  $A$  and  $B$  and renormalization, we get the message  $\lambda_D(B)$

$B$	$\lambda_D(B)$
+b	.5625
-b	.4375

b. We also use the equation from IBP algorithm to compute the family marginals

$\beta(+a, -b, +c) = \eta\phi_C(+c, +a, -b)\pi_C(+a)\pi_C(-b) = \eta 0.8 \times 0.6 \times 0.7 = \eta 0.336$ . Since this factor is naturally normalized,  $\eta = 1$ .

$\beta(+a, -b, +d) = \eta\phi_D(+d, +a, -b)\pi_D(+a)\pi_D(-b) = \eta 0.4 \times 0.5 \times 0.75 = \eta 0.15$ . This factor does not sum to one naturally since we have evidence on  $D$ . In this case  $\eta \approx 2.667$  and  $\beta(+a, -b, +d) = 0.4$ .

c. According to the IBP algorithm:

$\beta(A) = \eta\phi_A(A)\lambda_C(A)\lambda_D(A)$ . Therefore:

$A$	$\beta(A)$
+a	.6
-a	.4

Also,  $\beta(B) = \eta\phi_B(B)\lambda_C(B)\lambda_D(B)$ .

$B$	$\beta(B)$
+b	.3
-b	.7

d. We can compute  $\beta(A)$  by summing out variables  $B$  and  $C$  from  $\beta(A, B, C)$ :

$$\beta(A) = \sum_{B,C} \beta(A, B, C)$$

$A$	$B$	$C$	$\beta(A, B, C)$
+a	+b	+c	.162
+a	+b	-c	.018
+a	-b	+c	.336
+a	-b	-c	.084
-a	+b	+c	.084
-a	+b	-c	.036
-a	-b	+c	.168
-a	-b	-c	.112

Summing out  $C$ , we get:

$A$	$B$	$\beta(A, B)$
+a	+b	.18
+a	-b	.42
-a	+b	.12
-a	-b	.28

Summing out  $B$ , we get:

$A$	$\beta(A)$
+a	.6
-a	.4

We can also compute  $\beta(B)$  from  $\beta(A, B)$  eliminating  $A$

$B$	$\beta(B)$
+b	.3
-b	.7

We will now sum out variables from the family marginal  $\beta(A, B, D)$ . We can compute  $\beta(A)$  by summing out variables  $B$  and  $D$  from  $\beta(A, B, D)$ :

$A$	$B$	$D$	$\beta(A, B, D)$
+a	+b	+d	.2
+a	+b	-d	0
+a	-b	+d	.4
+a	-b	-d	0
-a	+b	+d	.1
-a	+b	-d	0
-a	-b	+d	.3
-a	-b	-d	0

Summing out  $D$ , we get:

$A$	$B$	$\beta(A, B)$
+a	+b	.2
+a	-b	.4
-a	+b	.1
-a	-b	.3

Summing out  $B$ , we get:

$A$	$\beta(A)$
+a	.6
-a	.4

We can also compute  $\beta(B)$  from  $\beta(A, B)$  eliminating  $A$

$B$	$\beta(B)$
+b	.3
-b	.7

- e. We did this as an intermediate step of the previous answer. We copy the answer above to facilitate reference.

$A$	$B$	$\beta(A, B)$ from $\beta(A, B, C)$
+a	+b	.18
+a	-b	.42
-a	+b	.12

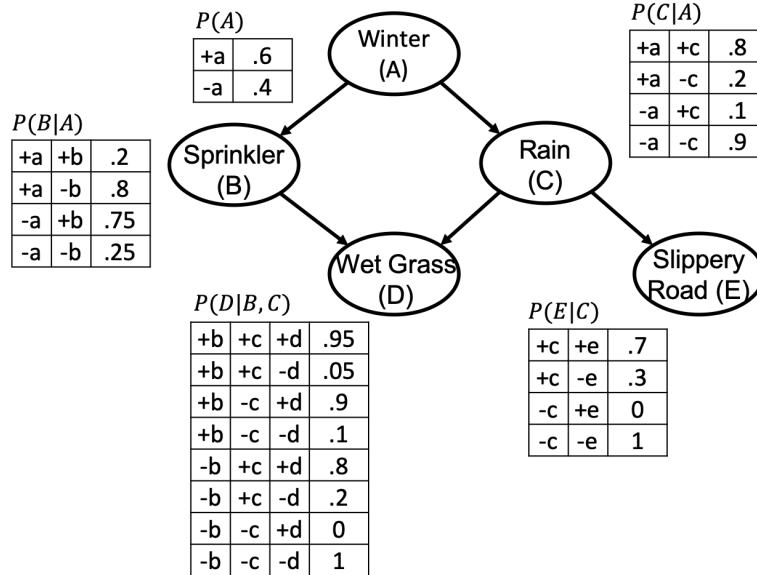
$A$	$B$	$\beta(A, B)$ from $\beta(A, B, C)$
-a	-b	.28

$A$	$B$	$\beta(A, B)$ from $\beta(A, B, D)$
+a	+b	.2
+a	-b	.4
-a	+b	.1
-a	-b	.3

- f. The marginals computed in (d) are consistent, while the ones in computed in (e) are not consistent.

## Question 2

Consider the following Bayesian network:



and the parameter  $\theta_{\bar{a}}$  representing the probability of  $A = false$  (i.e., it is not winter). For each of the following values of this parameter, 0.01, 0.4 and 0.99 do the following:

- Compute the probability of  $P(+d, +e)$ : wet grass and slippery road. You can use the code from previous tutorials.
- Estimate  $P(+d, +e)$  using forward sampling with sample sizes ranging from  $n = 100$  to  $n = 15,000$ .
- Generate a plot with  $n$  on the  $x$ -axis and the exact value of  $P(+d, +e)$  and the estimate for  $P(+d, +e)$  on the  $y$ -axis.
- Generate a plot with  $n$  on the  $x$ -axis and the exact variance of the estimate for  $P(+d, +e)$  and the sample variance on the  $y$ -axis.

### Answer

- We computed  $P(d, e)$  using our code from previous tutorials. We obtained the following results:  
 $P(+d, +e) = 0.4608$  for  $\theta_{\bar{a}} = 0.01$   $P(+d, +e) = 0.3044$  for  $\theta_{\bar{a}} = 0.4$   $P(+d, +e) = 0.0679$  for  $\theta_{\bar{a}} = 0.99$
- The following code will estimate  $P(+d, +e)$  using forward sampling:

```

# Declare `factors`: a factor dictionary for the Bayesian network parameters

theta_na_list = [.01, .4, .99]
true_p_d_e_list = [0.4608, 0.3044, 0.0679]
sample_sizes = range(100,15000,100)
for j, theta_na in enumerate(theta_na_list):
    p_d_e_list = []
    sample_var_list = []
    sample_mean_var_list = []
    factors['A']['table'][1] = 1 - theta_na
    for s in sample_sizes:
        d_e_count = 0
        for i in range(s):
            a = random() < factors['A']['table'][1]
            b = random() < factors['B']['table'][(a,1)]
            c = random() < factors['C']['table'][(a,1)]
            d = random() < factors['D']['table'][(b,c,1)]
            e = random() < factors['E']['table'][(c,1)]
            if d and e:
                d_e_count = d_e_count + 1

        p_d_e = d_e_count / s
        p_d_e_list.append(p_d_e)
        sample_var_list.append((sum([(1-p_d_e)**2 for i in range(d_e_count)]) + \
                                   sum([(0-p_d_e)**2 for i in range(s-d_e_count)])) / (s-1))
        sample_mean_var_list.append(true_p_d_e_list[j] * (1-true_p_d_e_list[j])/s)


plt.figure()
plt.axhline(y=true_p_d_e_list[j], color='r', linestyle='-')
plt.suptitle('P(a)='+str(1-theta_na))
plt.plot(range(100,15000,100), p_d_e_list, 'bo')
plt.xlabel("Sample size")
plt.ylabel("P(+d,+e)")
plt.show()

plt.figure()
plt.suptitle('P(a)='+str(1-theta_na))
plt.plot(range(100,15000,100), sample_var_list, 'b')
plt.xlabel("Sample size")
plt.ylabel("Sample variance")
plt.show()

plt.figure()
plt.suptitle('P(a)='+str(1-theta_na))
plt.plot(range(100,15000,100), sample_mean_var_list, 'b')
plt.xlabel("Sample size")
plt.ylabel("Sample mean variance")
plt.show()

```

c. You should get a plot like this, with the above code



- d. Due to the difference in scale, the code creates separated plots for sample variance and the exact variance of the estimate.

